

Stackable Certificate Programme in Artificial Intelligence

Cognitive Systems

NATURAL LANGUAGE COMPREHENSION AND PROCESSING

Fan Zhenzhen
zhenzhen@nus.edu.sg



© 2019 NUS. All rights reserved.

Page 1 of 34

Objective

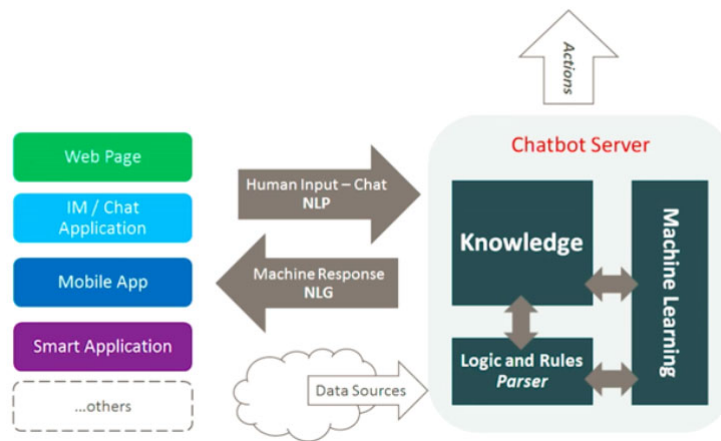
- To review and identify appropriate NLP techniques and solutions for Cognitive Systems to process natural language inputs
- Topics:
 - Recap...
 - The need of natural language comprehension
 - Natural language processing techniques



© 2019 NUS. All rights reserved.

Page 2 of 34

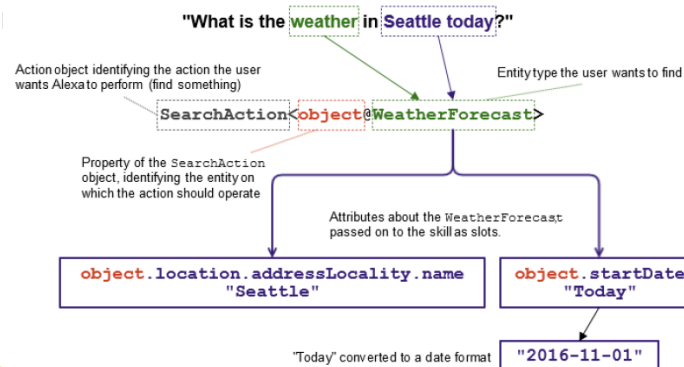
Recap...



<https://medium.com/analytics-vidhya/building-a-simple-chatbot-in-python-using-nltk-7c8c8215ac6e>

Recap...

- Natural language understanding is required for the system to understand the user's request
 - To map user's utterances to **intents**, which may
- CO



Recap...

What is the weather in Seattle today?

What's the weather in London?

Tell me the temperature now.

Intent Detection

Slots Detection

`SearchAction(WeatherForecast((Location 'Seattle'),(StartDate 'Today')))`

NLP Tasks Required

- Named Entity Recognition
- Entity Labelling
- Entity Linking
- Co-reference Resolution

Information Extraction

- Intent Detection (Classification)
- Slot Detection (NER)

What is information extraction?

- The automatic extraction of (possibly pre-specified) information from natural language documents
 - Facts about types of entities, events, relationships
- The automatic population of a structured information source (template, or logical form) from natural language documents
 - Documents may be semi-structured (eg., patents), unstructured (e.g., websites) or free text (e.g., documents)

Concept vs. Named Entity vs. Information

- **Name Entity** = lowest level of recognition by an IE system
 - Normally recognized by dictionaries or rules
- **Concept** = rule or heuristic to create an abstraction
 - Sometimes called a “natural class” = different people at different times and in different places would refer to the same referent with that concept
 - “president of the United States” vs. “president of the United Kingdom”
- **Information** = words, named entities, concepts which fulfill a need
 - So if you have a question, and a phrase answers that question, then that phrase is an example of information
 - Information is often regular, i.e., with a pattern
 - Eg, information about a person = name, age, sex, address, hp#, ...
 - Information about a company = name, address, stock symbol, Chairman, ...

What is information extraction?

- The automatic extraction of (possibly pre-specified) information from natural language documents
 - Facts about types of entities, events, relationships
- The automatic population of a structured information source (template) from natural language documents (i.e., create a table!)
 - Documents may be semi-structured (eg., patents), unstructured (e.g., websites) or free text (e.g., documents)

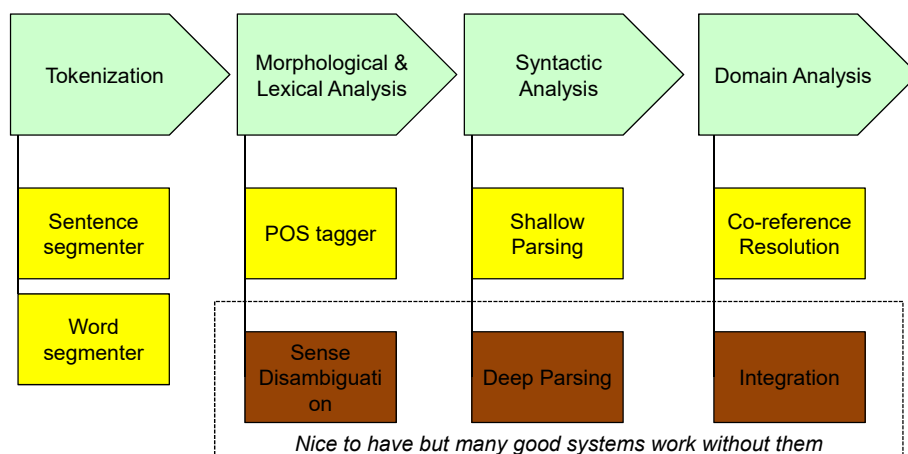
Types of IE systems

- Rule-based Systems
 - Hand-coded rules
 - Coded by linguists, with domain input
 - Iterative method based on document inspection
 - Slow but very good results
 - Induced (machine learning) rules
 - Fully machine learning
 - Given an annotated corpus, derive a basis set of rules that cover a pre-determined % of the annotated examples (and only the annotated examples)
 - Heuristic approach: one rule at a time!
 - Hybrid systems – machine learning to fine-tune the rules

Types of IE systems

- Statistics-based Systems
 - Start with a well-annotated corpus
 - Depending on the method (e.g., Hidden Markov Models), derive statistical rules to create a model that generates the examples
 - Advantages compared to Rule based systems
 - Language independent (within representational limits)
 - No linguistic or domain knowledge needed in the team
 - Relatively small effort in creating the models
 - Issues
 - The complexity moves to the corpus – must be well annotated and must cover the full space of possibilities
 - Requires very large number of training examples to get good results

Main components of an IE system



Tokenization

- To break a stream of characters into tokens

Great location with a little bit of history.



Great location with a little bit of history

- This is done by identifying token delimiters
 - Whitespace characters such as *space, tab, newline*
 - Punctuation characters like *() < > ! ? " ' "*
 - Other characters *., :- ' ' etc.*

Tokenization Challenges

- It seems simple, but...
 - ., :* between numbers are part of the number
 - .* can be part of an abbreviation or end of a sentence
 - '* can be a closing internal quote, indicate a possessive, or be part of another token

12.34

12,345

12:34

- .* can be part of an abbreviation or end of a sentence

U.S.A.

Dr.

- '* can be a closing internal quote, indicate a possessive, or be part of another token

My friend's

isn't

POS Tagging

- To determine POS or grammatical category of a term
 - Nouns, verbs, adjectives, adverbs, pronouns, determiners, prepositions, conjunctions, etc.
 - LDC Penn Tree Bank has 36 categories with detailed information, e.g.

| | | | |
|-----|--|-----|---------------------------------------|
| CC | Coordinating conjunction | UH | Interjection |
| CD | Cardinal number | VB | Verb, base form |
| DT | Determiner | VBD | Verb, past tense |
| EX | Existential <i>there</i> | VBG | Verb, gerund or present participle |
| FW | Foreign word | VBN | Verb, past participle |
| IN | Preposition or subordinating conjunction | VBP | Verb, non-3rd person singular present |
| JJ | Adjective | VBZ | Verb, 3rd person singular present |
| JJR | Adjective, comparative | WDT | Wh-determiner |
| JJS | Adjective, superlative | WP | Wh-pronoun |

POS Tagging

- Dictionary with word-POS correspondence is needed
- Challenge – POS disambiguation (words with >1 POS)
 - E.g. “book” can be a noun (“my book”) or a verb (“to book a room”)
- Example:
 - About six and a half hours later, Mr. Armstrong opened the landing craft's hatch, stepped slowly down the ladder and declared as he planted the first human footprint on the lunar crust: “That's one small step for man, one giant leap for mankind.”

IN/ About CD/ six CC/ and DT/ a JJ/ half NNS/ hours RB/ later ,/ , NNP/ Mr. NNP/ Armstrong VBD/ opened DT/ the NN/ landing NN/ craft POS/ 's NN/ hatch ,/ , VBD/ stepped RB/ slowly IN/ down DT/ the NN/ ladder CC/ and VBD/ declared IN/ as PRP/ he VBD/ planted DT/ the JJ/ first NN/ human NN/ footprint IN/ on DT/ the NN/ lunar NN/ crust :/ : ``/ " DT/ That VBZ/ 's CD/ one JJ/ small NN/ step IN/ for NN/ man ,/ , CD/ one JJ/ giant NN/ leap IN/ for NN/ mankind ./ . "/ "

Generated by UIUC POS Tagger

POS Taggers

- Rule-based - e.g. Brill's tagger by Eric Brill
 - Error-driven transformation-based tagger
 - Initially assign the most frequent tag to each word, based on dictionary and morphological rules
 - Contextual rules are then applied repeatedly to correct any errors
- Stochastic taggers – e.g. CLAWS, Viterbi, Baum-Welch, etc.
 - based on Hidden Markov Models (HMMs) and n-gram probabilities
 - Manually tagged corpus is needed to estimate probabilities
- Many machine learning methods have also been applied
- Stanford's Statistical NLP website lists many free taggers

Shallow Parsing / Chunking

- To identify phrases in a text (noun phrases, verb phrases, and prepositional phrases, etc.)
- Example:
 - About six and a half hours later, Mr. Armstrong opened the landing craft's hatch, stepped slowly down the ladder and declared as he planted the first human footprint on the lunar crust: "That's one small step for man, one giant leap for mankind."

[NP About six and a half hours] [ADVP later] , [NP Mr. Armstrong] [VP opened] [NP the landing craft] [NP 's hatch] , [VP stepped] [ADVP slowly] [PP down] [NP the ladder] and [VP declared] [SBAR as] [NP he] [VP planted] [NP the first human footprint] [PP on] [NP the lunar crust] : "[NP That] [VP 's] [NP one small step] [PP for] [NP man] , [NP one giant leap] [PP for] [NP mankind] ."

Generated by UIUC chunker

Shallow Parsing / Chunking

- After morphological analysis and disambiguation, using information of lemmata, morphological information, and word order configuration
- Largely stochastic techniques based on probabilities derived from an annotated corpus
- Avoiding the complexity of full parsing, faster, more robust
- Useful in Information Extraction, Summary Generation, and Question Answering

Name Entity Recognition

- Recognition of particular types of proper noun phrases, specifically persons, organizations, locations, and sometimes money, dates, times, and percentages.
- Very useful in text mining applications, by turning verbose text data into a more compact structural form

[LOC Houston], Monday, July 21 -- Men have landed and walked on the moon. Two [MISC Americans], astronauts of [ORG Apollo] 11, steered their fragile four-legged lunar module safely and smoothly to the historic landing yesterday at 4:17:40 P.M., Eastern daylight time. [PER Neil A. Armstrong], the 38-year-old civilian commander, radioed to earth and the mission control room here: "[LOC Houston], [ORG Tranquility Base] here; the Eagle has landed."

Generated by UIUC NER system

Rule-based NER Systems

- Rule-based systems can and do work well
 - Corpus is relatively static (in terms of vocabulary, language structure, etc.)
 - Can be fast especially in well-defined limited domains (compared to annotating training examples)
- A typical rule-based system comprises
 - Set of rules
 - Policies to control when and how (multiple) rules are applied, e.g., order, looping.

What does a rule look like?

- Lexical pattern matching
- Form:
 - Match(pattern) then Do(action)

```

Rule: Company1                                     from gate.ac.uk
( ( {Token.orthography == upperInitial} )+
  {Lookup.kind == companyDesignator}
):match
-->
:match.NamedEntity = { kind=company, rule="Company1" }
  
```

When to use statistics based systems?

- Many top performing systems are statistics based
 - Machine learning (ML) on very large corpora is state-of-the-art
- Annotation based corpora for training
 - You have a well annotated corpora with many features
 - Various ML techniques from simple to sophisticated
 - Relatively homogeneous real data (not training data) in any given domain. Note that models don't transfer well across domains
 - You don't have domain or language resources in that area

Simple model is at token level

- Text is a linear sequence of tokens (such as words)
- Token boundaries can be fairly easily derived in some languages, e.g., space & punctuation for English, but much harder for others, e.g., Chinese
- Simple tokenization
 - Dictionary based
 - Colocation frequencies (see next page)
- Alternatives
 - Ignore multi-unit tokens
 - Bi-grams, tri-grams, multi-grams

Popular models

- Hidden Markov Models (HMM)
 - Simple, joint probability
- Conditional Random Fields (CRF)
 - Conditional probability
 - Considers features of current token, and of preceding n tokens (window= n)
- Similarity algorithms
 - Measure distance of group of words to a dictionary list
 - Works especially well for jargon and other terminology
- Support Vector Machines (SVM)
 - Training method for standard perceptron
 - Optimize the points to determine the hyperplane dividing the positive training samples from the negative ones

What is coreference?

- Coreference resolution
 - Determine relationship between entities which are related
 - Identity relation (morning star vs. evening star)
 - Whole-part relation
 - Simple version
 - Determine entities which have the same referent
 - Anaphora (Pronouns)
 - Proper names, proper nouns, noun phrases,...
 - Definite descriptions (may be time dependent)
 - Usain Bolt & "the fastest man in the world"

Co-reference Examples

- Anaphora
 - The elephant stepped on the rabbit and it died.
 - The elephant stepped on the landmine and it died.
- Proper nouns
 - John Smith and Mary Brown were married this morning. The groom was dressed in a white tuxedo while the bride was...
- Definite descriptions
 - Usain Bolt has won the Olympic 100m gold medal. The fastest man in the world successfully defended his title last night.

Intent Classification

- Using labelled data to build machine learning models that can classify input into intent classes (supervised learning)
- SVM/NB/LR/DT/KNN
- Pre-process the input text into features (vector model)



| | amazing | service | lost | glamour | disappoint | brilliant | super | expensive | noisy | ... |
|------|---------|---------|------|---------|------------|-----------|-------|-----------|-------|-----|
| Doc1 | 1 | 1 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | |
| Doc2 | 0 | 0 | 1 | 1 | 1 | 0 | 0 | 1 | 0 | |
| Doc3 | 0 | 0 | 0 | 1 | 0 | 0 | 1 | 0 | 0 | |
| Doc4 | 0 | 0 | 0 | 0 | 2 | 0 | 0 | 1 | 1 | |
| ... | | | | | | | | | | |

Typical Pre-processing

- Tokenization
- Case normalisation (to lowercase)
- Lemmatization/stemming
 - To reduce the words to its root form
 - E.g. classes -> class, ran -> run , production -> produce
- Punctuation removal
- Stopword removal
 - To remove extremely common words (with little meaning) like functional words (the, a, of...)

Indexing

- Many text mining applications are based on vector representation of documents (term-document matrix) using “bag-of-words” approach

$$\begin{pmatrix}
 & T_1 & T_2 & \dots & T_t \\
 D_1 & w_{11} & w_{21} & \dots & w_{t1} \\
 D_2 & w_{12} & w_{22} & \dots & w_{t2} \\
 \vdots & \vdots & \vdots & & \vdots \\
 \vdots & \vdots & \vdots & & \vdots \\
 D_n & w_{1n} & w_{2n} & \dots & w_{tn}
 \end{pmatrix}$$

T : term
 D : document
 w : weight of the term

- Usually only content words (adjectives, adverbs, nouns, and verbs) are used as vector features.

Term Weighting

- Binary
 - 0 or 1, simply indicating whether a word has occurred in the document (but that's not very helpful).
- Frequency-based
 - *term frequency*, the frequency of words in the document, which provides additional information that can be used to contrast with other documents.

| | amazing | service | lost | glamour | disappoint | brilliant | super | expensive | noisy | ... |
|------|---------|---------|------|---------|------------|-----------|-------|-----------|-------|-----|
| Doc1 | 1 | 1 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | |
| Doc2 | 0 | 0 | 1 | 1 | 1 | 0 | 0 | 1 | 0 | |
| Doc3 | 0 | 0 | 0 | 1 | 0 | 0 | 1 | 0 | 0 | |
| Doc4 | 0 | 0 | 0 | 0 | 2 | 0 | 0 | 1 | 1 | |
| ... | | | | | | | | | | |

tf-idf Indexing

- To modify the frequency of a word in a document by the perceived importance of the word (the *inverse document frequency*), widely used in information retrieval
 - When a word appears in many documents, it's considered unimportant.
 - When the word is relatively unique and appears in few documents, it's important.

$$tf-idf_{t,d} = tf_{t,d} * idf_t$$

$$idf_t = \log \frac{N}{df_t}$$

- $tf_{t,d}$: term frequency – number of occurrences of term t in document d
- idf_t : inverted document frequency of term t

N : the total number of documents in the corpus

df_t : the document frequency of term t , i.e., the number of documents that contain the term.

tf-idf Indexing – An Example

| TERM VECTOR MODEL BASED ON $w_i = tf_i \cdot IDF_i$ | | | | | | | | | | | |
|---|----------------|----------------|----------------|----------------|-----------------|-------------------|-----------------------------------|--------|----------------|----------------|----------------|
| Query, Q: "gold silver truck" | | | | | | | | | | | |
| D ₁ : "Shipment of gold damaged in a fire" | | | | | | | | | | | |
| D ₂ : "Delivery of silver arrived in a silver truck" | | | | | | | | | | | |
| D ₃ : "Shipment of gold arrived in a truck" | | | | | | | | | | | |
| D = 3; IDF = log(D/df _i) | | | | | | | | | | | |
| Terms | Counts, tf_i | | | | | | Weights, $w_i = tf_i \cdot IDF_i$ | | | | |
| | Q | D ₁ | D ₂ | D ₃ | df _i | D/df _i | IDF _i | Q | D ₁ | D ₂ | D ₃ |
| a | 0 | 1 | 1 | 1 | 3 | 3/3 = 1 | 0 | 0 | 0 | 0 | 0 |
| arrived | 0 | 0 | 1 | 1 | 2 | 3/2 = 1.5 | 0.1761 | 0 | 0 | 0.1761 | 0.1761 |
| damaged | 0 | 1 | 0 | 0 | 1 | 3/1 = 3 | 0.4771 | 0 | 0.4771 | 0 | 0 |
| delivery | 0 | 0 | 1 | 0 | 1 | 3/1 = 3 | 0.4771 | 0 | 0 | 0.4771 | 0 |
| fire | 0 | 1 | 0 | 0 | 1 | 3/1 = 3 | 0.4771 | 0 | 0.4771 | 0 | 0 |
| gold | 1 | 1 | 0 | 1 | 2 | 3/2 = 1.5 | 0.1761 | 0.1761 | 0.1761 | 0 | 0.1761 |
| in | 0 | 1 | 1 | 1 | 3 | 3/3 = 1 | 0 | 0 | 0 | 0 | 0 |
| of | 0 | 1 | 1 | 1 | 3 | 3/3 = 1 | 0 | 0 | 0 | 0 | 0 |
| silver | 1 | 0 | 2 | 0 | 1 | 3/1 = 3 | 0.4771 | 0.4771 | 0 | 0.9542 | 0 |
| shipment | 0 | 1 | 0 | 1 | 2 | 3/2 = 1.5 | 0.1761 | 0 | 0.1761 | 0 | 0.1761 |
| truck | 1 | 0 | 1 | 1 | 2 | 3/2 = 1.5 | 0.1761 | 0.1761 | 0 | 0.1761 | 0.1761 |

Note that in this example, stopwords and very common words are not removed, and terms are not reduced to root terms.

<http://www.milsila.com/term-vector/term-vector-3.html>

Cosine Similarity

- A similarity measure between two vectors (input and candidate response)
- by measuring the cosine of the angle between them

$$Sim(D_i, D_j) = \frac{D_i \cdot D_j}{|D_i| \cdot |D_j|} = \frac{\sum_k w_{ki} w_{kj}}{\sqrt{\sum_k w_{ki}^2} \sqrt{\sum_k w_{kj}^2}}$$

- Example: Given 3 document vectors shown here

$$|D_1| = \sqrt{0.1761^2 + 0.4771^2 + 0.1761^2} = \sqrt{0.2896} = 0.5382$$

$$|D_2| = \sqrt{0.4771^2 + 0.4771^2 + 0.1761^2 + 0.1761^2} = \sqrt{0.5173} = 0.7192$$

$$|D_3| = \sqrt{0.1761^2 + 0.4771^2 + 0.9542^2 + 0.1761^2} = \sqrt{1.2001} = 1.0955$$

$$Sim(D_1, D_2) = (0.1761 \cdot 0.1761) / (0.5382 \cdot 0.7192) = 0.0801$$

$$Sim(D_1, D_3) = (0.4771 \cdot 0.9542 + 0.1761 \cdot 0.1761) / (0.5382 \cdot 1.0955) = 0.8246$$

| D ₁ | D ₂ | D ₃ |
|----------------|----------------|----------------|
| 0 | 0 | 0 |
| 0 | 0 | 0.1761 |
| 0 | 0.4771 | 0 |
| 0 | 0 | 0.4771 |
| 0 | 0.4771 | 0 |
| 0.1761 | 0.1761 | 0 |
| 0 | 0 | 0 |
| 0 | 0 | 0 |
| 0.4771 | 0 | 0.9542 |
| 0 | 0.1761 | 0 |
| 0.1761 | 0 | 0.1761 |